

Unchecked Exceptions – The Controversy

Unchecked Exceptions – The Controversy

Because the Java programming language does not require methods to catch or to specify unchecked exceptions ([RuntimeException](#), [Error](#), and their subclasses), programmers may be tempted to write code that throws only unchecked exceptions or to make all their exception subclasses inherit from [RuntimeException](#). Both of these shortcuts allow programmers to write code without bothering with compiler errors and without bothering to specify or to catch any exceptions. Although this may seem convenient to the programmer, it sidesteps the intent of the catch or specify requirement and can cause problems for others using your classes.

Why did the designers decide to force a method to specify all uncaught checked exceptions that can be thrown within its scope? Any [Exception](#) that can be thrown by a method is part of the method's public programming interface. Those who call a method must know about the exceptions that a method can throw so that they can decide what to do about them. These exceptions are as much a part of that method's programming interface as its parameters and return value.

The next question might be: "If it's so good to document a method's API, including the exceptions it can throw, why not specify runtime exceptions too?"

Runtime exceptions represent problems that are the result of a programming problem, and as such, the API client code cannot reasonably be expected to recover from them or to handle them in any way. Such problems include arithmetic exceptions, such as dividing by zero; pointer exceptions, such as trying to access an object through a null reference; and indexing exceptions, such as attempting to access an array element through an index that is too large or too small.

Runtime exceptions can occur anywhere in a program, and in a typical one they can be very numerous. Having to add runtime exceptions in every method declaration would reduce a program's clarity. Thus, the compiler does not require that you catch or specify runtime exceptions (although you can).

One case where it is common practice to throw a [RuntimeException](#) is when the user calls a method incorrectly. For example, a method can check if one of its arguments is incorrectly null. If an argument is null, the method might throw a [NullPointerException](#), which is an unchecked exception.

Generally speaking, do not throw a [RuntimeException](#) or create a subclass of [RuntimeException](#) simply because you don't want to be bothered with specifying the exceptions your methods can throw.

Here's the bottom line guideline: If a client can reasonably be expected to recover from an exception, make it a checked exception. If a client cannot do anything to recover from the exception, make it an unchecked exception.

Advantages of Exceptions

Now that you know what exceptions are and how to use them, it's time to learn the advantages of using exceptions in your programs.

Advantage 1: Separating Error-Handling Code from "Regular" Code

Exceptions provide the means to separate the details of what to do when something out of the ordinary happens from the main logic of a program. In traditional programming, error detection, reporting, and handling often lead to confusing spaghetti code. For example, consider the pseudocode method here that reads an entire file into memory.

```
1  readFile {
2      open the file;
3      determine its size;
4      allocate that much memory;
5      read the file into memory;
6      close the file;
7  }
```

At first glance, this function seems simple enough, but it ignores all the following potential errors.

- What happens if the file can't be opened?
- What happens if the length of the file can't be determined?
- What happens if enough memory can't be allocated?
- What happens if the read fails?
- What happens if the file can't be closed?

To handle such cases, the `readFile` function must have more code to do error detection, reporting, and handling. Here is an example of what the function might look like.

```
1  errorCodeType readFile {
2      initialize errorCode = 0;
3
4      open the file;
5      if (theFileIsOpen) {
6          determine the length of the file;
7          if (gotTheFileLength) {
8              allocate that much memory;
9              if (gotEnoughMemory) {
10
```

```

10         read the file into memory;
11         if (readFailed) {
12             errorCode = -1;
13         }
14     } else {
15         errorCode = -2;
16     }
17 } else {
18     errorCode = -3;
19 }
20 close the file;
21 if (theFileDidntClose && errorCode == 0) {
22     errorCode = -4;
23 } else {
24     errorCode = errorCode and -4;
25 }
26 } else {
27     errorCode = -5;
28 }
29 return errorCode;
30 }

```

There's so much error detection, reporting, and returning here that the original seven lines of code are lost in the clutter. Worse yet, the logical flow of the code has also been lost, thus making it difficult to tell whether the code is doing the right thing: Is the file really being closed if the function fails to allocate enough memory? It's even more difficult to ensure that the code continues to do the right thing when you modify the method three months after writing it. Many programmers solve this problem by simply ignoring it — errors are reported when their programs crash.

Exceptions enable you to write the main flow of your code and to deal with the exceptional cases elsewhere. If the `readFile` function used exceptions instead of traditional error-management techniques, it would look more like the following.

```

1  readFile {
2      try {
3          open the file;

```

```
4         determine its size;
5         allocate that much memory;
6         read the file into memory;
7         close the file;
8     } catch (fileOpenFailed) {
9         doSomething;
10    } catch (sizeDeterminationFailed) {
11        doSomething;
12    } catch (memoryAllocationFailed) {
13        doSomething;
14    } catch (readFailed) {
15        doSomething;
16    } catch (fileCloseFailed) {
17        doSomething;
18    }
19 }
```

Note that exceptions don't spare you the effort of doing the work of detecting, reporting, and handling errors, but they do help you organize the work more effectively.

Advantage 2: Propagating Errors Up the Call Stack

A second advantage of exceptions is the ability to propagate error reporting up the call stack of methods. Suppose that the `readFile` method is the fourth method in a series of nested method calls made by the main program: `method1` calls `method2`, which calls `method3`, which finally calls `readFile`.

```
1  method1 {
2      call method2;
3  }
4
5  method2 {
6      call method3;
7  }
8
9  method3 {
10     call readFile;
11 }
```

Suppose also that `method1` is the only method interested in the errors that might occur within `readFile`. Traditional error-notification techniques force `method2` and `method3` to propagate the error codes returned by `readFile` up the call stack until the error codes finally reach `method1`—the only method that is interested in them.

```
1  method1 {
2      errorCodeType error;
3      error = call method2;
4      if (error)
5          doErrorProcessing;
6      else
7          proceed;
8  }
9
10 errorCodeType method2 {
11     errorCodeType error;
12     error = call method3;
13     if (error)
14         return error;
15     else
16         proceed;
17 }
18
19 errorCodeType method3 {
20     errorCodeType error;
21     error = call readFile;
22     if (error)
23         return error;
24     else
25         proceed;
26 }
```

Recall that the Java runtime environment searches backward through the call stack to find any methods that are interested in handling a particular exception. A method can duck any exceptions thrown within it, thereby allowing a method farther up the call stack to catch it. Hence, only the methods that care about errors have to worry about detecting errors.

```
1 | method1 {
2 |     try {
3 |         call method2;
4 |     } catch (exception e) {
5 |         doErrorProcessing;
6 |     }
7 | }
8 |
9 | method2 throws exception {
10 |     call method3;
11 | }
12 |
13 | method3 throws exception {
14 |     call readFile;
15 | }
```

However, as the pseudocode shows, ducking an exception requires some effort on the part of the middleman methods. Any checked exceptions that can be thrown within a method must be specified in its **throws** clause.

Advantage 3: Grouping and Differentiating Error Types

Because all exceptions thrown within a program are objects, the grouping or categorizing of exceptions is a natural outcome of the class hierarchy. An example of a group of related exception classes in the Java platform are those defined in [java.io](#) – [IOException](#) and its descendants. [IOException](#) is the most general and represents any type of error that can occur when performing I/O. Its descendants represent more specific errors. For example, [FileNotFoundException](#) means that a file could not be located on disk.

A method can write specific handlers that can handle a very specific exception. The [FileNotFoundException](#) class has no descendants, so the following handler can handle only one type of exception.

```
1 | catch (FileNotFoundException e) {
2 |     ...
3 | }
```

A method can catch an exception based on its group or general type by specifying any of the exception's superclasses in the catch statement. For example, to catch all I/O exceptions, regardless of their specific type, an exception handler specifies an [IOException](#) argument.

```
1 | catch (IOException e) {  
2 |     ...  
3 | }
```

This handler will be able to catch all I/O exceptions, including [FileNotFoundException](#), [EOFException](#), and so on. You can find details about what occurred by querying the argument passed to the exception handler. For example, use the following to print the stack trace.

```
1 | catch (IOException e) {  
2 |     // Output goes to System.err.  
3 |     e.printStackTrace();  
4 |     // Send trace to stdout.  
5 |     e.printStackTrace(System.out);  
6 | }
```

You could even set up an exception handler that handles any [Exception](#) with the handler here.

```
1 | // A (too) general exception handler  
2 | catch (Exception e) {  
3 |     ...  
4 | }
```

The [Exception](#) class is close to the top of the [Throwable](#) class hierarchy. Therefore, this handler will catch many other exceptions in addition to those that the handler is intended to catch. You may want to handle exceptions this way if all you want your program to do, for example, is print out an error message for the user and then exit.

In most situations, however, you want exception handlers to be as specific as possible. The reason is that the first thing a handler must do is determine what

type of exception occurred before it can decide on the best recovery strategy. In effect, by not catching specific errors, the handler must accommodate any possibility. Exception handlers that are too general can make code more error-prone by catching and handling exceptions that weren't anticipated by the programmer and for which the handler was not intended.

As noted, you can create groups of exceptions and handle exceptions in a general fashion, or you can use the specific exception type to differentiate exceptions and handle exceptions in an exact fashion.

Summary

A program can use exceptions to indicate that an error occurred. To throw an exception, use the `throw` statement and provide it with an exception object — a descendant of [Throwable](#) — to provide information about the specific error that occurred. A method that throws an uncaught, checked exception must include a `throws` clause in its declaration.

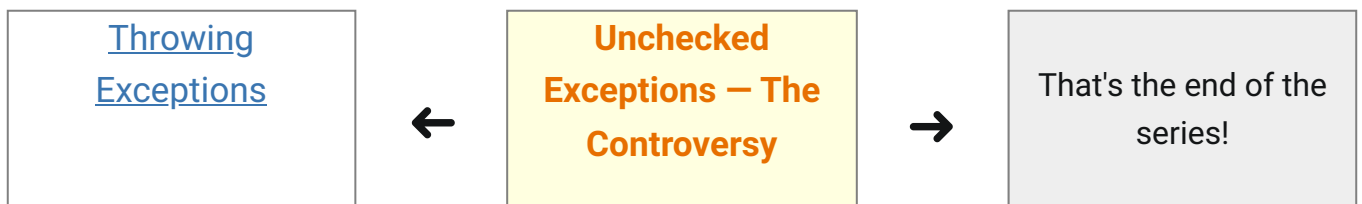
A program can catch exceptions by using a combination of the `try`, `catch`, and `finally` blocks.

- The `try` block identifies a block of code in which an exception can occur.
- The `catch` block identifies a block of code, known as an exception handler, that can handle a particular type of exception.
- The `finally` block identifies a block of code that is guaranteed to execute, and is the right place to close files, recover resources, and otherwise clean up after the code enclosed in the `try` block.

The `try` statement should contain at least one `catch` block or a `finally` block and may have multiple `catch` blocks.

The class of the exception object indicates the type of exception thrown. The exception object can contain further information about the error, including an error message. With exception chaining, an exception can point to the exception that caused it, which can in turn point to the exception that caused it, and so on.

Last update: September 14, 2021



[Home](#) > [Tutorials](#) > [Exceptions](#) > Unchecked Exceptions – The Controversy